



CS10 Head TA
Michael Ball

The Beauty and Joy of Computing

Lecture #8 Concurrency



"KOOMEY'S LAW" – EFFICIENCY 2X EVERY 18 MO

Prof Jonathan Koomey looked at 6 decades of data and found that energy efficiency of computers doubles roughly every 18 months. This is even more relevant as battery-powered devices become more popular. Restated, it says that for a fixed computing load, the amount of battery you need drops by half every 18 months. This was true before transistors!



www.technologyreview.com/computing/38548/



Concurrency: A Definition

Concurrency: A property of computer systems in which several computations are executing simultaneously, and potentially interacting with each other.





Concurrency is Everywhere!

Examples:

- Mouse cursor movement while Snap! calculates.
- Screen clock advances while typing in a text.
- Busy cursor spins while browser connects to server, waiting for response
- Walking while chewing gum





Concurrency & Parallelism

Intra-computer

- Today's lecture
- Multiple computing "helpers" are cores within one machine
- Aka "multi-core"
 - Although GPU parallelism is also "intra-computer"



HD & C800.com

Inter-computer


- Future lecture
- Multiple computing "helpers" are different machines
- Aka "distributed computing"
 - Grid & cluster computing



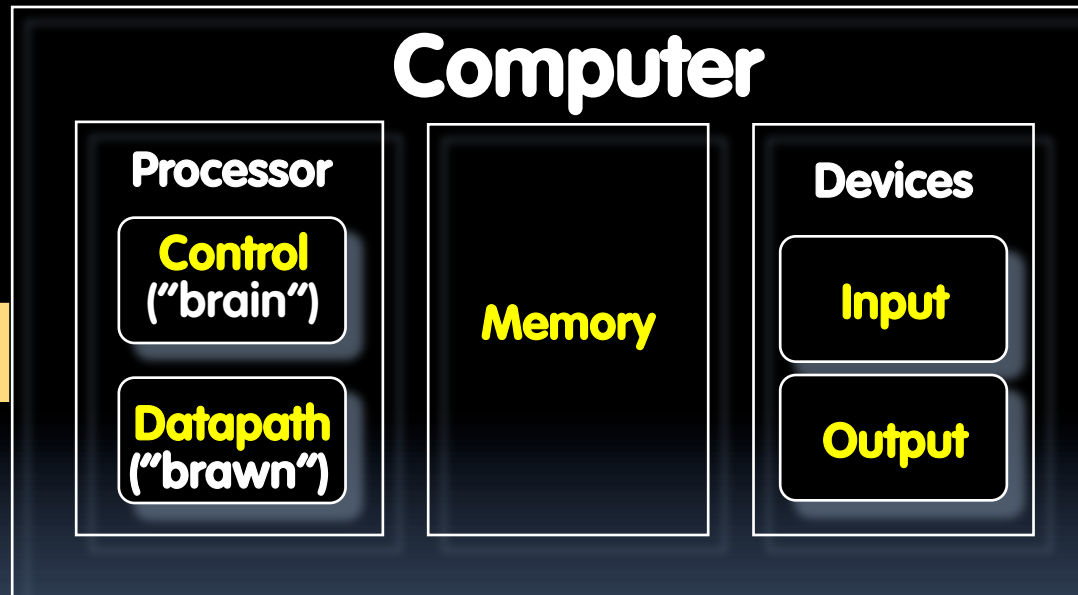


Anatomy: 5 components of any Computer

John von Neumann
invented this
architecture



- a) Control
- b) Datapath
- c) Memory
- d) Input
- e) Output

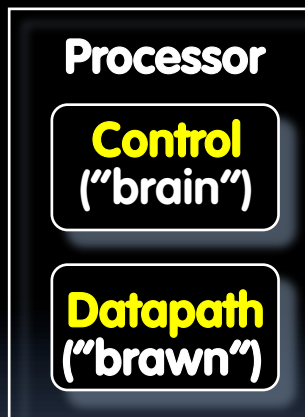


What causes the most headaches
for SW and HW designers with
multi-core computing?



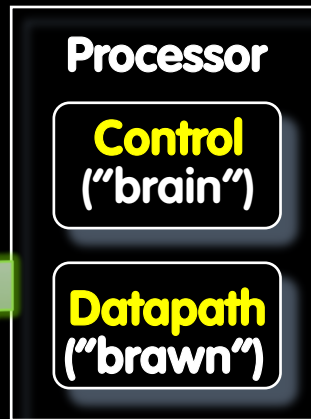


But what is INSIDE a Processor?

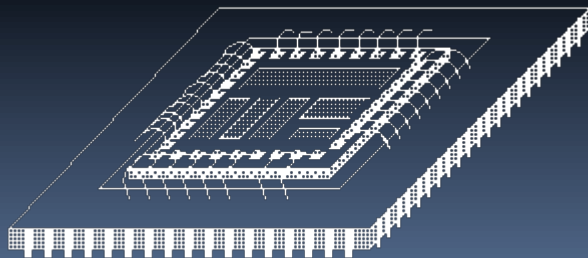




But what is INSIDE a Processor?



Bare Processor Die



Chip in Package

- Primarily Crystalline Silicon
- 1 mm – 25 mm on a side
- 2009 “feature size” (aka process)
~ 45 nm = 45×10^{-9} m
(then 32, 22, and 16 [by yr 2013])
- **100 - 1000M transistors**
- 3 - 10 conductive layers
- “CMOS” (complementary metal oxide semiconductor) - most common
- Package provides:
 - spreading of chip-level signal paths to board-level
 - heat dissipation.
- Ceramic or plastic with gold wires.


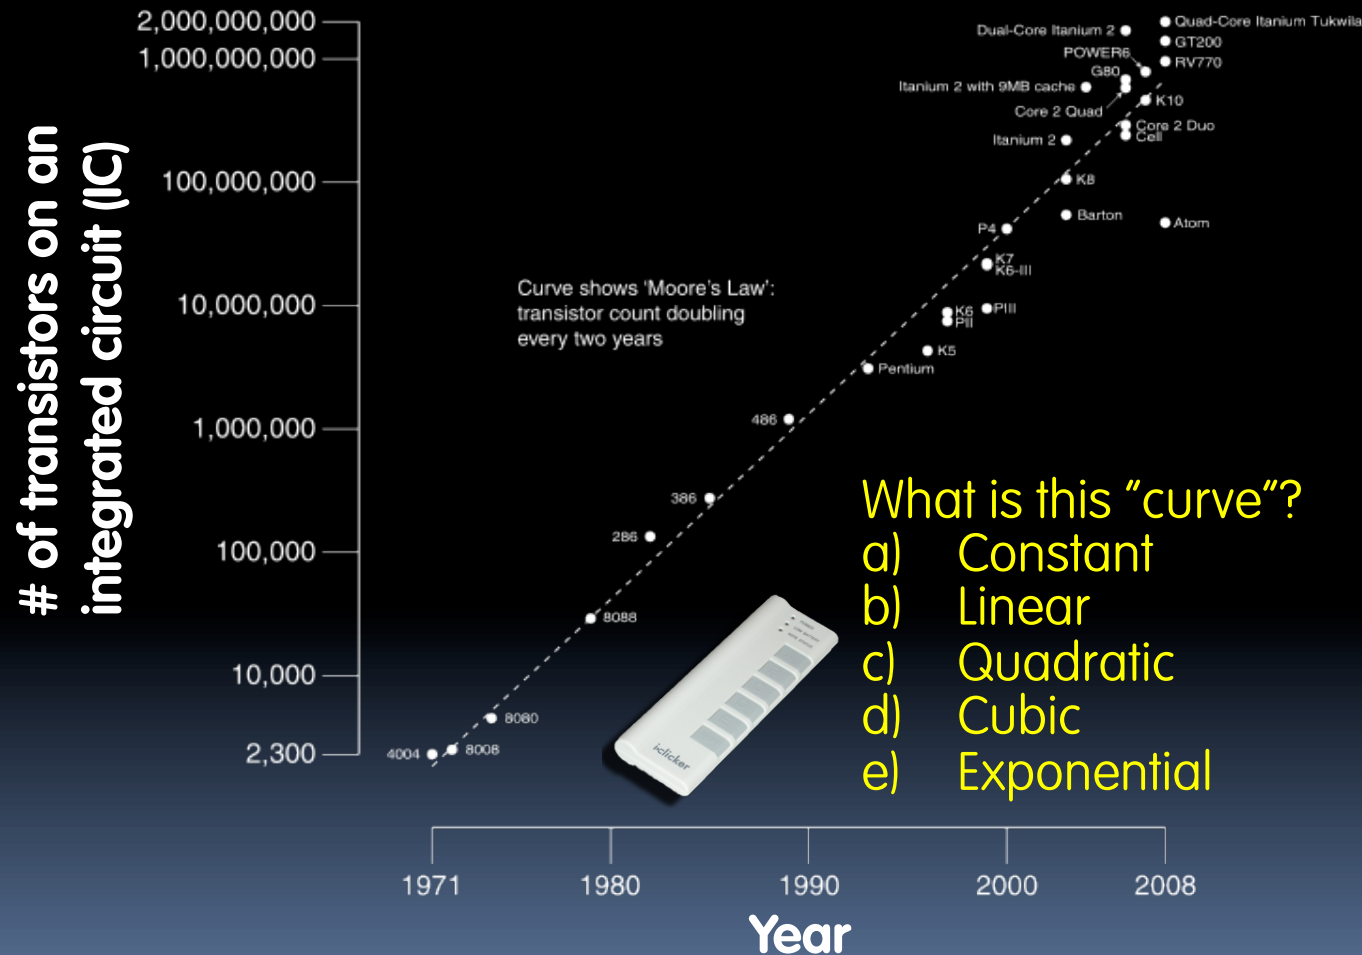




en.wikipedia.org/wiki/Moore's_law

Moore's Law

Predicts: 2X Transistors / chip every 2 years

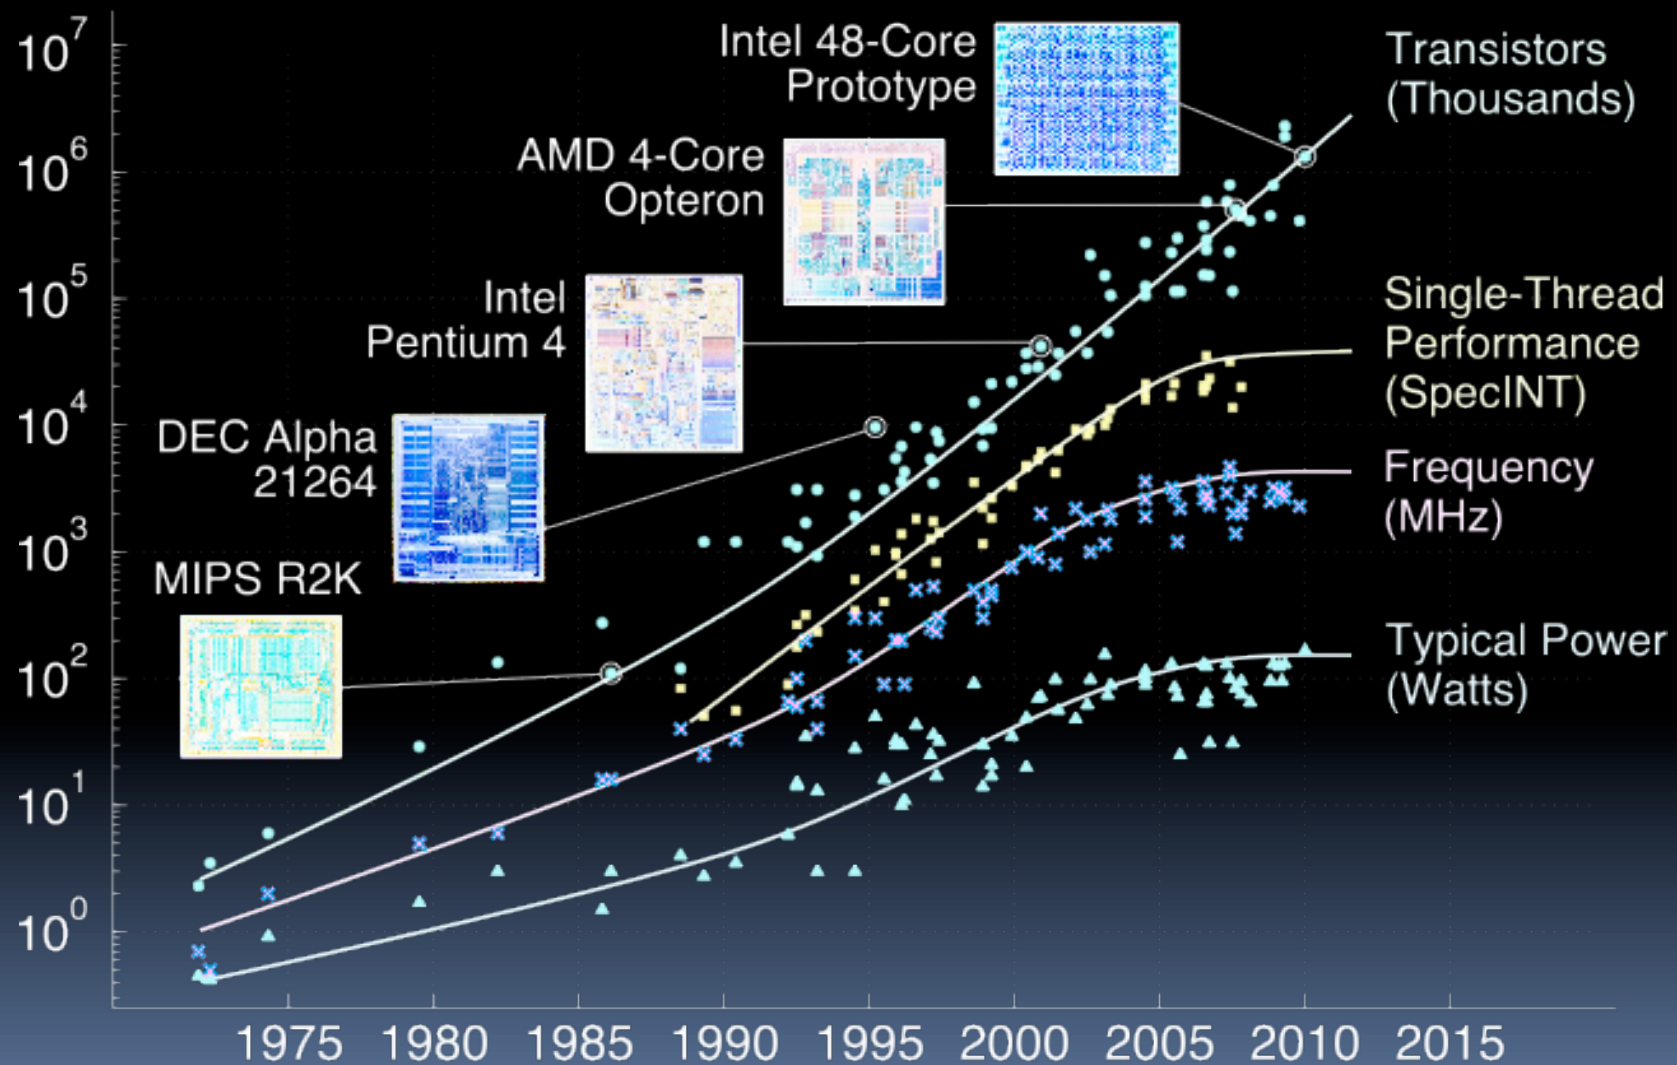


Gordon Moore
Intel Cofounder
B.S. Cal 1950!





Moore's Law and related curves

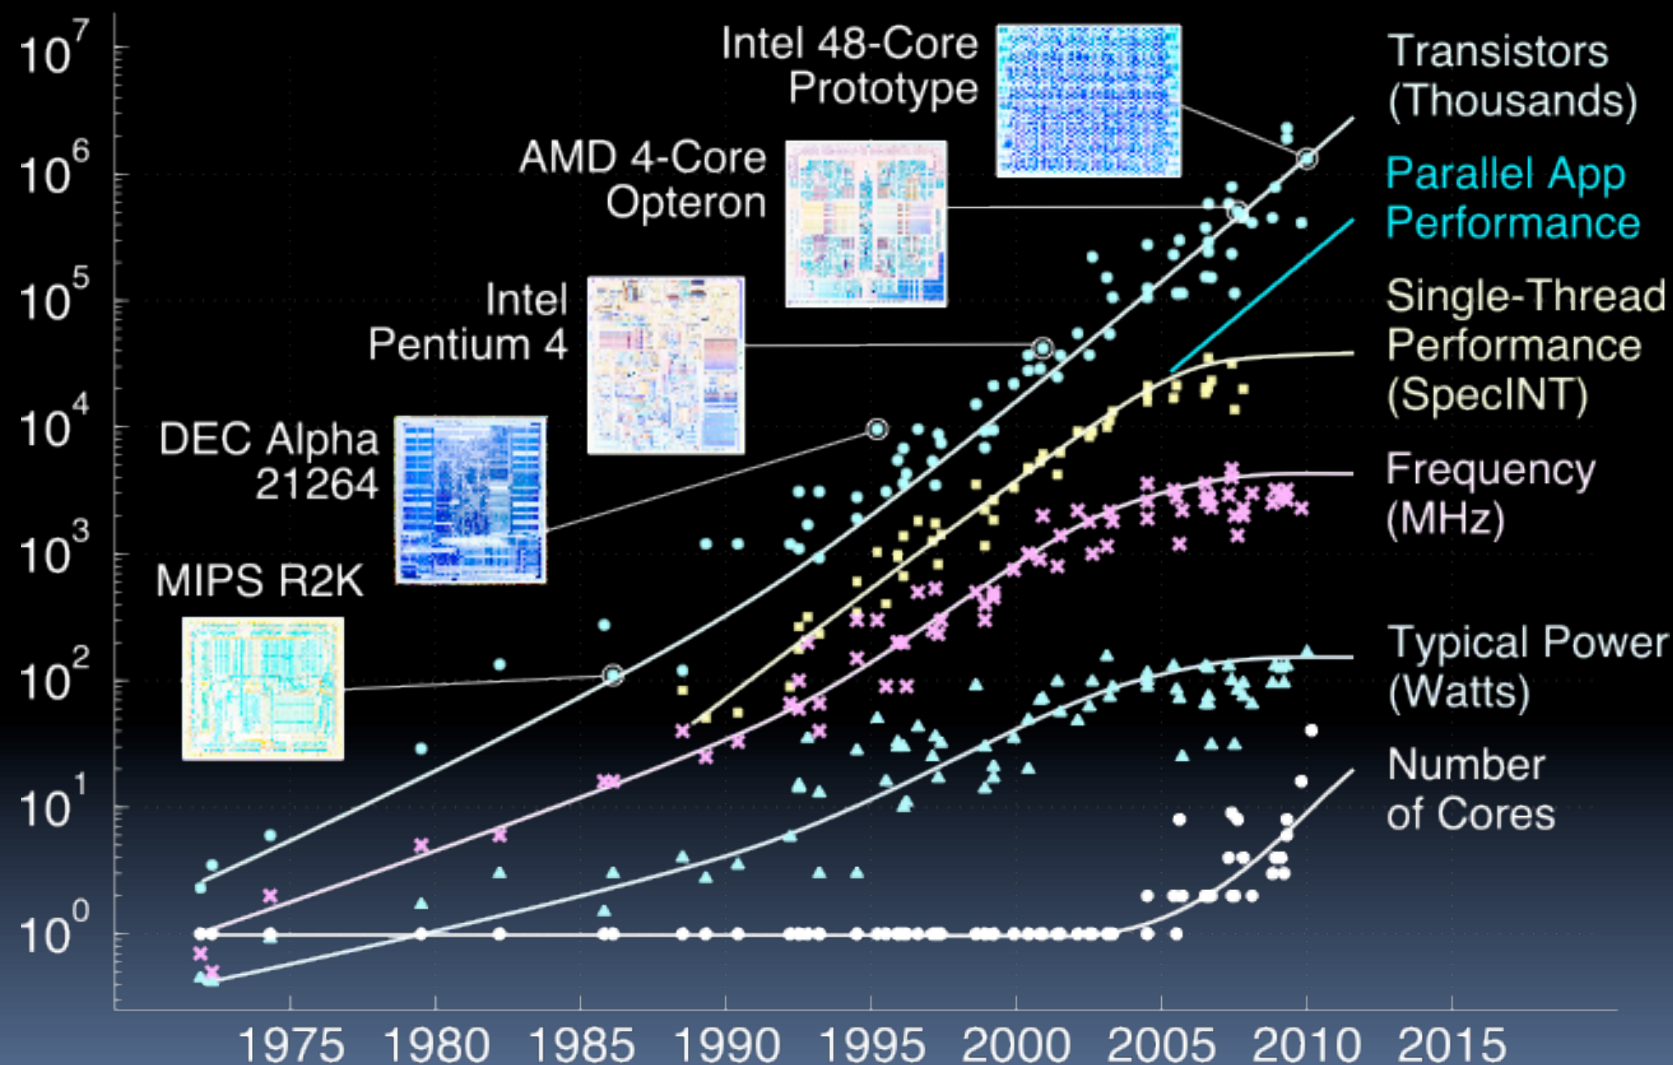


Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond





Moore's Law and related curves




Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond





Power Density Prediction circa 2000



Source: S. Borkar (Intel)

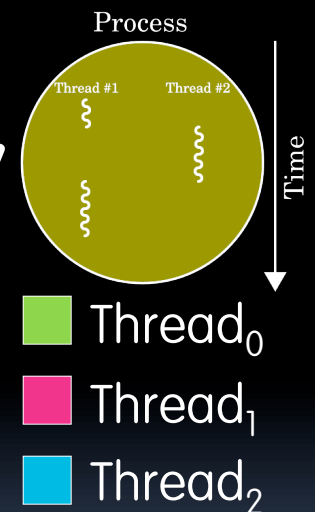




Background: Threads

- A **Thread** stands for “thread of execution”, is a single stream of instructions
 - A program / process can **split**, or **fork** itself into separate threads, which can (in theory) execute simultaneously.
 - An easy way to describe/think about parallelism

- A single CPU can execute many threads by **Time Division Multiplexing**



- **Multithreading** is running multiple threads through the same hardware

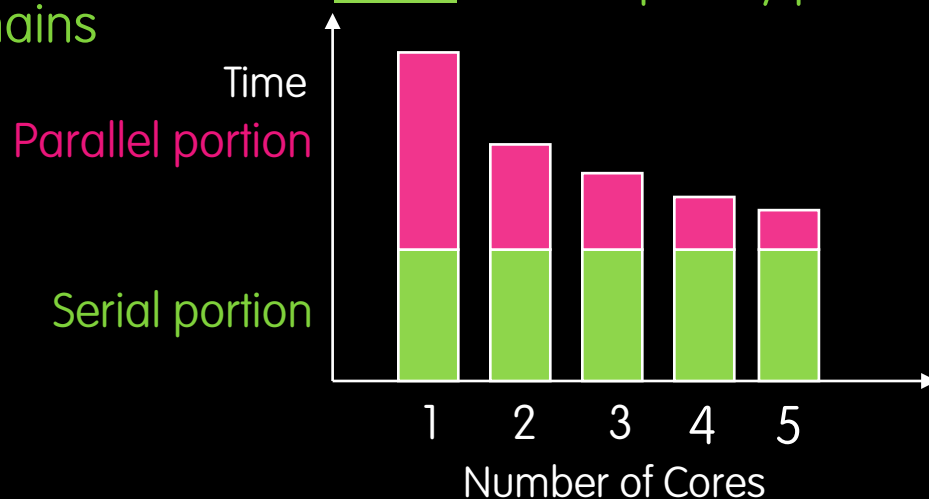




en.wikipedia.org/wiki/Amdahl's_law

Speedup Issues : Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains



- s is serial fraction of program, P is # of cores (was processors)

- **Amdahl's law:**

$$\text{Speedup}(P) = \text{Time}(1) / \text{Time}(P)$$

$$\leq 1 / (s + [(1-s) / P]), \text{ and as } P \rightarrow \infty$$

$$\leq 1 / s$$

- Even if the parallel portion of your application speeds up perfectly, **your performance may be limited by the sequential portion**





Speedup Issues : Overhead

- Even assuming no sequential portion, there's...
 - Time to think how to **divide the problem up**
 - Time to **hand out** small “work units” to workers
 - All workers may **not work equally fast**
 - Some **workers may fail**
 - There may be **contention for shared resources**
 - Workers could **overwriting each others' answers**
 - You may have to **wait until the last worker returns** to proceed (the slowest / weakest link problem)
 - There's **time to put the data back together** in a way that looks as if it were done by one





Life in a multi-core world...


- This “sea change” to multi-core parallelism means that the computing community has to rethink:
 - a) Languages
 - b) Architectures
 - c) Algorithms
 - d) Data Structures
 - e) All of the above





But parallel programming is hard!

- What if two people were calling withdraw at the same time?
 - E.g., balance=100 and two withdraw 75 each
 - Can anyone see what the problem *could* be?
 - This is a **race condition**
- In most languages, this is a problem.
 - In Snap!, the system doesn't let two of these run at once.





"Non-Deterministic" Parallel Code

- Two (or more) scripts are running at the same time, BUT we don't know what order they will be run in!
- Each individual script runs its blocks in order, but the processor (Snap!) will swap between running script A and script B.





How Many Possible Outputs?

- We want this code to draw a cute winky-face, but there's a problem with parallelizing it!
- How many possible outputs can we have?
- A) 1
- B) 3
- C) 4
- D) 7
- E) 8





Another concurrency problem ... **deadlock!**

- Two people need to draw a graph but there is only one pencil and one ruler.
 - One grabs the pencil
 - One grabs the ruler
 - Neither release what they hold, waiting for the other to release
- **Livelock** also possible
 - Movement, no progress





Summary

- “Sea change” of computing because of inability to cool CPUs means we’re now in multi-core world
- This brave new world offers lots of potential for innovation by computing professionals, but challenges persist

