


The Beauty and Joy of Computing


Lecture #8 Concurrency



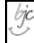
CS10 Head TA
Michael Ball

"KOOMEY'S LAW" – EFFICIENCY 2X EVERY 18 MO

Prof Jonathan Koomey looked at 6 decades of data and found that energy efficiency of computers doubles roughly every 18 months. This is even more relevant as battery-powered devices become more popular. Restated, it says that for a fixed computing load, the amount of battery you need drops by half every 18 months. This was true before transistors!




www.technologyreview.com/computing/38548/



Concurrency: A Definition

Concurrency: A property of computer systems in which several computations are executing simultaneously, and potentially interacting with each other.

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (2)

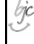


Concurrency is Everywhere!

Examples:



- Mouse cursor movement while Snap! calculates.
- Screen clock advances while typing in a text.
- Busy cursor spins while browser connects to server, waiting for response
- Walking while chewing gum

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (3)




Concurrency & Parallelism


Intra-computer	Inter-computer
<ul style="list-style-type: none"> ▪ Today's lecture ▪ Multiple computing "helpers" are cores <u>within one machine</u> ▪ Aka "multi-core" <ul style="list-style-type: none"> ▫ Although GPU parallism is also "intra-computer" 	<ul style="list-style-type: none"> ▪ Future lecture ▪ Multiple computing "helpers" are <u>different machines</u> ▪ Aka "distributed computing" <ul style="list-style-type: none"> ▫ Grid & cluster computing


UC Berkeley "The Beauty and Joy of Computing" : Concurrency (4)




Anatomy: 5 components of any Computer



John von Neumann
invented this
architecture






Computer

Processor Control ("brain") Datapath ("brawn")	Memory	Devices Input Output
---	---------------	-----------------------------------

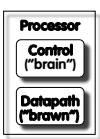
a) Control
b) Datapath
c) Memory
d) Input
e) Output

What causes the most headaches for SW and HW designers with multi-core computing?

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (5)



But what is INSIDE a Processor?



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (6)

But what is INSIDE a Processor?

- Primarily Crystalline Silicon
- 1 mm – 25 mm on a side
- 2009 “feature size” (aka process) ~ 45 nm = 45×10^{-9} m (then 32, 22, and 16 [by yr 2013])
- 100 - 1000M transistors
- 3 - 10 conductive layers
- “CMOS” (complementary metal oxide semiconductor) - most common
- Package provides:
 - spreading of chip-level signal paths to board-level
 - heat dissipation.
- Ceramic or plastic with gold wires.

UC Berkeley “The Beauty and Joy of Computing” : Concurrency (7)

Moore's Law

en.wikipedia.org/wiki/Moore's_law

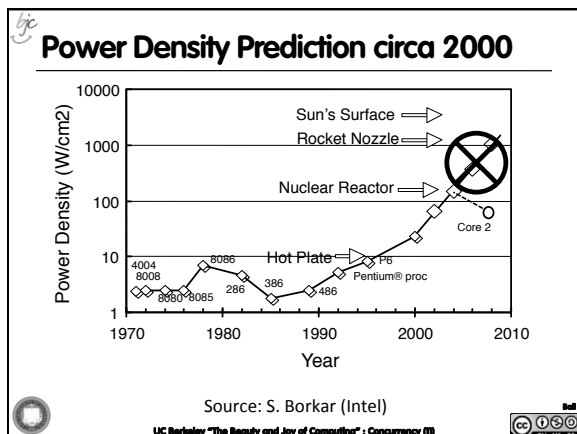
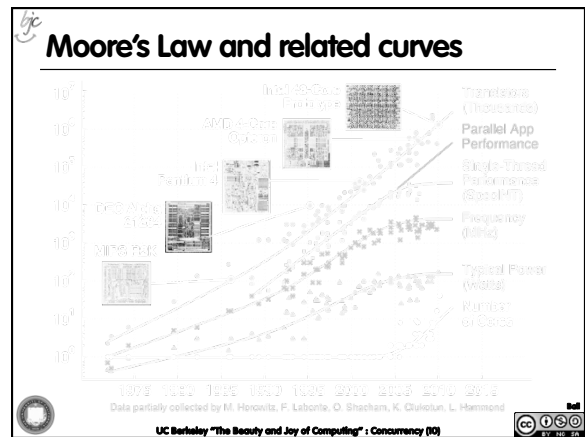
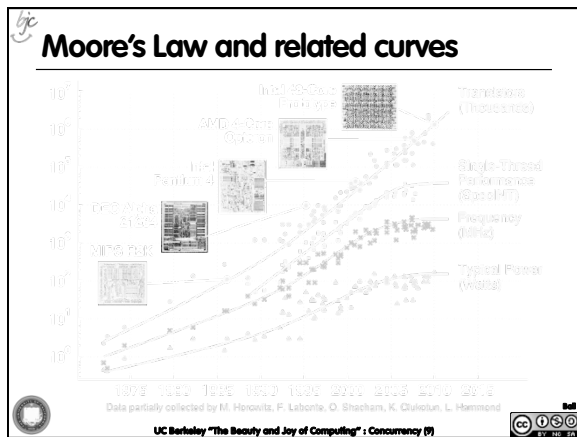
Predicts: 2X Transistors / chip every 2 years

What is this “curve”?

- Constant
- Linear
- Quadratic
- Cubic
- Exponential

Gordon Moore
Intel Co-founder
B.S. Cal 1950!

UC Berkeley “The Beauty and Joy of Computing” : Concurrency (8)



Background: Threads

- A **Thread** stands for “thread of execution”, is a single stream of instructions
 - A program / process can split, or fork itself into separate threads, which can (in theory) execute simultaneously.
 - An easy way to describe/think about parallelism
- A single CPU can execute many threads by **Time Division Multiplexing**
- **Multithreading** is running multiple threads through the same hardware

UC Berkeley “The Beauty and Joy of Computing” : Concurrency (12)

en.wikipedia.org/wiki/Amdahl's_law

Speedup Issues : Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains

Time
Parallel portion
Serial portion
Number of Cores

- s is serial fraction of program, P is # of cores (was processors)
- Amdahl's law:**

$$\text{Speedup}(P) = \text{Time}(1) / \text{Time}(P)$$

$$\leq 1 / (s + (1-s) / P), \text{ and as } P \rightarrow \infty$$

$$\leq 1 / s$$
- Even if the parallel portion of your application speeds up perfectly, your performance may be limited by the sequential portion

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (13)

Speedup Issues : Overhead

- Even assuming no sequential portion, there's...
 - Time to think how to divide the problem up
 - Time to hand out small "work units" to workers
 - All workers may not work equally fast
 - Some workers may fail
 - There may be contention for shared resources
 - Workers could overwriting each others' answers
 - You may have to wait until the last worker returns to proceed (the slowest / weakest link problem)
 - There's time to put the data back together in a way that looks as if it were done by one

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (14)

Life in a multi-core world...

- This "sea change" to multi-core parallelism means that the computing community has to rethink:
 - Languages
 - Architectures
 - Algorithms
 - Data Structures
 - All of the above

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (15)

But parallel programming is hard!

- What if two people were calling withdraw at the same time?
 - E.g., balance=100 and two withdraw 75 each
 - Can anyone see what the problem *could* be?
 - This is a race condition
- In most languages, this is a problem.
 - In Snap!, the system doesn't let two of these run at once.

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (16)

"Non-Deterministic" Parallel Code


- Two (or more) scripts are running at the same time, BUT we don't know what order they will be run in!
- Each individual script runs its blocks in order, but the processor (Snap!) will swap between running script A and script B.

UC Berkeley "The Beauty and Joy of Computing" : Concurrency (17)

How Many Possible Outputs?


- We want this code to draw a cute winky-face, but there's a problem with parallelizing it!
- How many possible outputs can we have?
 - A) 1
 - B) 3
 - C) 4
 - D) 7
 - E) 8




UC Berkeley "The Beauty and Joy of Computing" : Concurrency (18)

 en.wikipedia.org/wiki/Deadlock

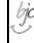
Another concurrency problem ... deadlock!

- Two people need to draw a graph but there is only one pencil and one ruler.
 - One grabs the pencil
 - One grabs the ruler
 - Neither release what they hold, waiting for the other to release
- Livelock also possible
 - Movement, no progress



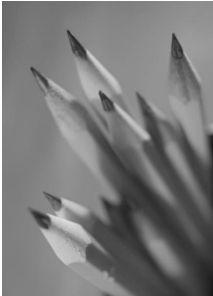
  




UC Berkeley "The Beauty and Joy of Computing" : Concurrency (19)



Summary

- "Sea change" of computing because of inability to cool CPUs means we're now in multi-core world
- This brave new world offers lots of potential for innovation by computing professionals, but challenges persist



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (20)